Student Assessment Risk Prediction using Logistic Regression from Scratch

Course: Data Science and Machine Learning Practicals
Dataset: Open University Learning Analytics Dataset — `assessments.csv`
Algorithm: Logistic Regression (implemented from scratch — no ML libraries)
Objective: Classify each student assessment as **High Risk (1)** or **Low Risk (0)**

Introduction

Educational institutions generate vast amounts of data through their Learning Management Systems (LMS), Student Information Systems (SIS), and assessment platforms. **Educational Data Mining (EDM)** is the discipline of applying data analysis and machine learning techniques to this data in order to extract actionable insights about student learning behaviour, engagement, and performance.

A critical challenge in higher education is **identifying assessments that carry a disproportionately high risk of poor student outcomes** — either because they are weighted heavily, scheduled late in the semester when students are already under stress, or because they are formal examinations. When institutions can flag these high-risk assessments early, academic advisors and faculty can **intervene proactively** — adjusting support resources, scheduling revision sessions, or providing additional guidance before performance deteriorates.

Predictive analytics transforms raw assessment metadata into an early warning signal. Rather than waiting for exam results to identify struggling students, a predictive model scores each assessment at the start of the semester, enabling a targeted and efficient support strategy. This notebook builds such a model using **Logistic Regression implemented entirely from scratch**, without relying on any machine learning library, demonstrating the mathematical foundations of binary classification in a real educational context.

Literature Survey

The application of machine learning to educational data has a rich research history. The five key works summarised below directly inform the design of this notebook.

Romero & Ventura (2010) — *Predicting Student Performance Using Data Mining*, Computers & Education is the foundational reference for this project. The authors applied classification algorithms including Decision Trees and Neural Networks to Moodle LMS logs. They computationally defined the *at-risk student* as one with attendance below 75% **and** assignment submission rate below 60%, finding

an 82% dropout likelihood for this group. This work established that assessment timing, weight, and engagement metrics are the most predictive features — a principle directly applied in our feature engineering pipeline.

Arnold & Pistilli (2012) — *Early Alert Systems Using Predictive Analytics*, Purdue University describes a production deployment across 32,000 students over five semesters. The authors used **Logistic Regression** on GPA trend, attendance rate, and LMS login frequency, achieving an AUC of 0.79 for dropout prediction. A dashboard-driven intervention reduced at-risk dropout rates by 21%. This paper directly motivates our choice of Logistic Regression as the model algorithm and validates the use of assessment weight and type as proxy features when direct student marks are unavailable.

Baker & Siemens (2014) — *Learning Analytics and Educational Data Mining* distinguishes between Learning Analytics (institution-level dashboards) and EDM (automated pattern discovery). They propose a layered architecture — raw event store → aggregated warehouse → ML feature store → visualisation layer — which this notebook emulates at a small scale through its preprocessing and feature engineering pipeline.

Chaudhuri & Dayal (EDU-DW literature) — *A Data Warehouse Framework for Educational Institutions* proposes a multi-subject-area data warehouse with conformed dimensions across enrollment, academic, and financial marts. The paper highlights the difficulty of integrating heterogeneous sources (SIS, LMS, HR), which is why we work with a single, clean assessments file and engineer derived features to compensate for missing student-level records.

Kimball Group — *Dimensional Modeling for Higher Education* applies the Kimball Bus Architecture to a university setting. The grain of the central fact table is defined as one row per student per course per term — the same grain as the OULAD assessments dataset, where each row represents one assessment event within a module presentation.

Across all five papers, a consistent set of predictive signals emerges: **assessment weight, assessment type (especially formal examinations), temporal placement within the semester (date), and module identity**. These signals form the basis of the feature set used in this notebook.

Importing Libraries

In [1]:
```
import pandas as pd
import numpy as np
import math
import random
import matplotlib.pyplot as plt
from collections import Counter

# Reproducibility — fix seeds so results are consistent across runs
```

```
random.seed(42)
np.random.seed(42)

print("All libraries imported successfully")
print("Note: No scikit-learn, tensorflow, or any ML library is used.")
print("      Logistic Regression is implemented entirely from scratch.")
```

```
All libraries imported successfully
Note: No scikit-learn, tensorflow, or any ML library is used.
      Logistic Regression is implemented entirely from scratch.
```

Dataset Loading and Understanding

The dataset used in this notebook is the **Open University Learning Analytics Dataset (OULAD)** `assessments.csv` file. OULAD is a publicly available dataset released by The Open University (UK) covering 32,593 students across 22 courses and 7 semesters. The `assessments.csv` file contains metadata for every assessment event — it does **not** contain individual student scores, making feature engineering from the available columns essential.

In [2]:
```
# Load the assessments dataset
df = pd.read_csv('assessments.csv')
print("Dataset loaded successfully")
df.head(10)
```

```
Dataset loaded successfully
```

Out[2]:

	code_module	code_presentation	id_assessment	assessment_type	date	weig
0	AAA	2013J	1752	TMA	19.0	1(
1	AAA	2013J	1753	TMA	54.0	2(
2	AAA	2013J	1754	TMA	117.0	2(
3	AAA	2013J	1755	TMA	166.0	2(
4	AAA	2013J	1756	TMA	215.0	3(
5	AAA	2013J	1757	Exam	NaN	10(
6	AAA	2014J	1758	TMA	19.0	1(
7	AAA	2014J	1759	TMA	54.0	2(
8	AAA	2014J	1760	TMA	117.0	2(
9	AAA	2014J	1761	TMA	166.0	2(

In [3]:
```
# Shape of the dataset
print("Shape:", df.shape)
print(f"  Rows    : {df.shape[0]} assessments")
print(f"  Columns : {df.shape[1]} attributes")
```

```
Shape: (206, 6)
  Rows    : 206 assessments
  Columns : 6 attributes
```

In [4]:
```
# Column names
print("Columns:")
```

```
print(list(df.columns))
```

```
Columns:
['code_module', 'code_presentation', 'id_assessment', 'assessment_type',
'date', 'weight']
```

In [5]:
```
# Check for missing values
print("Missing values per column:")
print(df.isnull().sum())
```

```
Missing values per column:
code_module           0
code_presentation     0
id_assessment         0
assessment_type       0
date                 11
weight                0
dtype: int64
```

In [6]:
```
# Summary statistics for numeric columns
df[['date', 'weight']].describe().round(2)
```

Out[6]:

	date	weight
count	195.00	206.00
mean	145.01	20.87
std	76.00	30.38
min	12.00	0.00
25%	71.00	0.00
50%	152.00	12.50
75%	222.00	24.25
max	261.00	100.00

Column Descriptions

Column	Type	Description
`code_module`	Categorical	Anonymised module identifier (e.g. AAA, BBB, CCC ...)
`code_presentation`	Categorical	Semester/year identifier (e.g. 2013J, 2014B)
`id_assessment`	Integer	Unique identifier for each assessment event
`assessment_type`	Categorical	Type of assessment — TMA (Tutor-Marked Assignment), CMA (Computer-Marked Assignment), or Exam
`date`	Numeric	Day number within the semester when the assessment is due (relative to module start)
`weight`	Numeric	Percentage weight of this assessment in the final module grade

The dataset contains **no direct target variable** (no pass/fail or risk label). We will engineer a binary `risk_level` target in the preprocessing section.

Data Preprocessing

Since all ML libraries are excluded, every preprocessing step — encoding, normalisation, feature engineering, and target creation — is implemented manually using only `pandas` and `numpy`.

In [7]:
```
# ── Step 1: Handle missing values ──────────────────────────────────
# The 'date' column may contain NaN values (assessments with no fixed due
# We fill NaN dates with the median date — a conservative neutral imputat

df['date'] = df['date'].fillna(df['date'].median())
print("Missing values after imputation:")
print(df.isnull().sum())
```

```
Missing values after imputation:
code_module          0
code_presentation    0
id_assessment        0
assessment_type      0
date                 0
weight               0
dtype: int64
```

In [8]:
```
# ── Step 2: Manual encoding of categorical columns ─────────────────
# We build dictionaries mapping each category string to an integer ID.
# This avoids any library-based label encoder.

# Encode code_module
module_categories = sorted(df['code_module'].unique())
module_map = {m: i for i, m in enumerate(module_categories)}
df['module_id'] = df['code_module'].map(module_map)

# Encode code_presentation
pres_categories = sorted(df['code_presentation'].unique())
pres_map = {p: i for i, p in enumerate(pres_categories)}
df['presentation_id'] = df['code_presentation'].map(pres_map)

# Encode assessment_type
type_categories = sorted(df['assessment_type'].unique())
type_map = {t: i for i, t in enumerate(type_categories)}
df['assessment_type_id'] = df['assessment_type'].map(type_map)

print("Module mapping     :", module_map)
print("Presentation mapping:", pres_map)
print("Assessment type mapping:", type_map)
```

```
Module mapping     : {'AAA': 0, 'BBB': 1, 'CCC': 2, 'DDD': 3, 'EEE': 4, 'F
FF': 5, 'GGG': 6}
Presentation mapping: {'2013B': 0, '2013J': 1, '2014B': 2, '2014J': 3}
Assessment type mapping: {'CMA': 0, 'Exam': 1, 'TMA': 2}
```

In [9]:
```
# ── Step 3: Derived feature engineering ────────────────────────────
# We create four additional features that capture educational risk signal
```

```
# identified in the literature (Romero & Ventura 2010, Arnold & Pistilli

# is_exam: 1 if the assessment is a formal examination, 0 otherwise
df['is_exam'] = (df['assessment_type'] == 'Exam').astype(int)

# is_high_weight: 1 if the assessment carries >= 50% of the total module
df['is_high_weight'] = (df['weight'] >= 50).astype(int)

# assessment_period: divides the semester into three bands
#   Early (0-100 days), Mid (101-200 days), Late (201+ days)
def get_period(d):
    if d <= 100:
        return 0   # Early
    elif d <= 200:
        return 1   # Mid
    else:
        return 2   # Late

df['assessment_period'] = df['date'].apply(get_period)

# module_difficulty_score: average weight of all assessments in the same
# Higher average weight → more demanding module → higher difficulty signa
module_avg_weight = df.groupby('code_module')['weight'].transform('mean')
w_min = module_avg_weight.min()
w_max = module_avg_weight.max()
df['module_difficulty_score'] = (module_avg_weight - w_min) / (w_max - w_

print("Derived features created:")
print("  is_exam, is_high_weight, assessment_period, module_difficulty_sc
```

```
Derived features created:
  is_exam, is_high_weight, assessment_period, module_difficulty_score
```

In [10]:

```
# ── Step 4: Create the target variable 'risk_level' ──────────────────
# Risk is a derived binary label — there is no ground-truth risk column i
# Based on the literature and domain knowledge, an assessment is High Ris
#   (a) Its weight is >= 50% of the module grade, OR
#   (b) It is a formal Exam, OR
#   (c) It is due very late in the semester (day > 200) with non-trivial

def compute_risk(row):
    if row['weight'] >= 50:
        return 1
    if row['assessment_type'] == 'Exam':
        return 1
    if row['date'] > 200 and row['weight'] >= 25:
        return 1
    return 0

df['risk_level'] = df.apply(compute_risk, axis=1)

print("Target variable 'risk_level' created")
print()
print("Risk Level Distribution:")
vc = df['risk_level'].value_counts()
print(f"  Low Risk  (0) : {vc.get(0, 0)} assessments")
print(f"  High Risk (1) : {vc.get(1, 0)} assessments")
print(f"  Class balance : {vc.get(1,0)/len(df)*100:.1f}% high risk")
```

```
Target variable 'risk_level' created

Risk Level Distribution:
  Low Risk  (0) : 179 assessments
  High Risk (1) : 27 assessments
  Class balance : 13.1% high risk
```

In [11]:
```python
# ── Step 5: Normalise date and weight columns ──────────────────────────
# Min-max normalisation scales both columns to [0, 1] range.
# This ensures no single feature dominates gradient descent due to scale

def minmax_normalize(series):
    mn = series.min()
    mx = series.max()
    return (series - mn) / (mx - mn + 1e-9)

df['date_norm']   = minmax_normalize(df['date'])
df['weight_norm'] = minmax_normalize(df['weight'])

print("Normalisation applied to 'date' and 'weight'")
print(f"  date_norm   range: [{df['date_norm'].min():.4f}, {df['date_norm
print(f"  weight_norm range: [{df['weight_norm'].min():.4f}, {df['weight_
```

```
Normalisation applied to 'date' and 'weight'
  date_norm   range: [0.0000, 1.0000]
  weight_norm range: [0.0000, 1.0000]
```

In [12]:
```python
# ── Step 6: Build final feature matrix ─────────────────────────────────
# Select only the engineered numeric features for model training.

feature_cols = [
    'module_id', 'presentation_id', 'assessment_type_id',
    'date_norm', 'weight_norm',
    'is_exam', 'is_high_weight', 'assessment_period',
    'module_difficulty_score'
]

target_col = 'risk_level'

print("Feature columns selected:")
for i, f in enumerate(feature_cols, 1):
    print(f"  {i}. {f}")
print(f"\nTotal features : {len(feature_cols)}")
print(f"Total samples  : {len(df)}")
```

```
Feature columns selected:
  1. module_id
  2. presentation_id
  3. assessment_type_id
  4. date_norm
  5. weight_norm
  6. is_exam
  7. is_high_weight
  8. assessment_period
  9. module_difficulty_score

Total features : 9
Total samples  : 206
```

In [13]:
```
# Preview the processed dataset
df[feature_cols + [target_col]].head(10)
```

Out[13]:

	module_id	presentation_id	assessment_type_id	date_norm	weight_norm	is_e
0	0	1	2	0.028112	0.1	
1	0	1	2	0.168675	0.2	
2	0	1	2	0.421687	0.2	
3	0	1	2	0.618474	0.2	
4	0	1	2	0.815261	0.3	
5	0	1	1	0.562249	1.0	
6	0	3	2	0.028112	0.1	
7	0	3	2	0.168675	0.2	
8	0	3	2	0.421687	0.2	
9	0	3	2	0.618474	0.2	

Exploratory Data Analysis

Before modelling, we visualise the dataset to understand the distribution of assessments across types, modules, and risk levels. Understanding class balance and feature distributions helps us anticipate model behaviour.

In [14]:
```
# ── Plot 1: Assessment type distribution ──────────────────────────
type_counts = df['assessment_type'].value_counts()
labels  = type_counts.index.tolist()
counts  = type_counts.values.tolist()
colors  = ['#3498db', '#2ecc71', '#e74c3c']

plt.figure(figsize=(6, 4))
bars = plt.bar(labels, counts, color=colors[:len(labels)], edgecolor='bla
for bar, c in zip(bars, counts):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
             str(c), ha='center', fontsize=11, fontweight='bold')
plt.title('Assessment Type Distribution')
plt.ylabel('Number of Assessments')
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

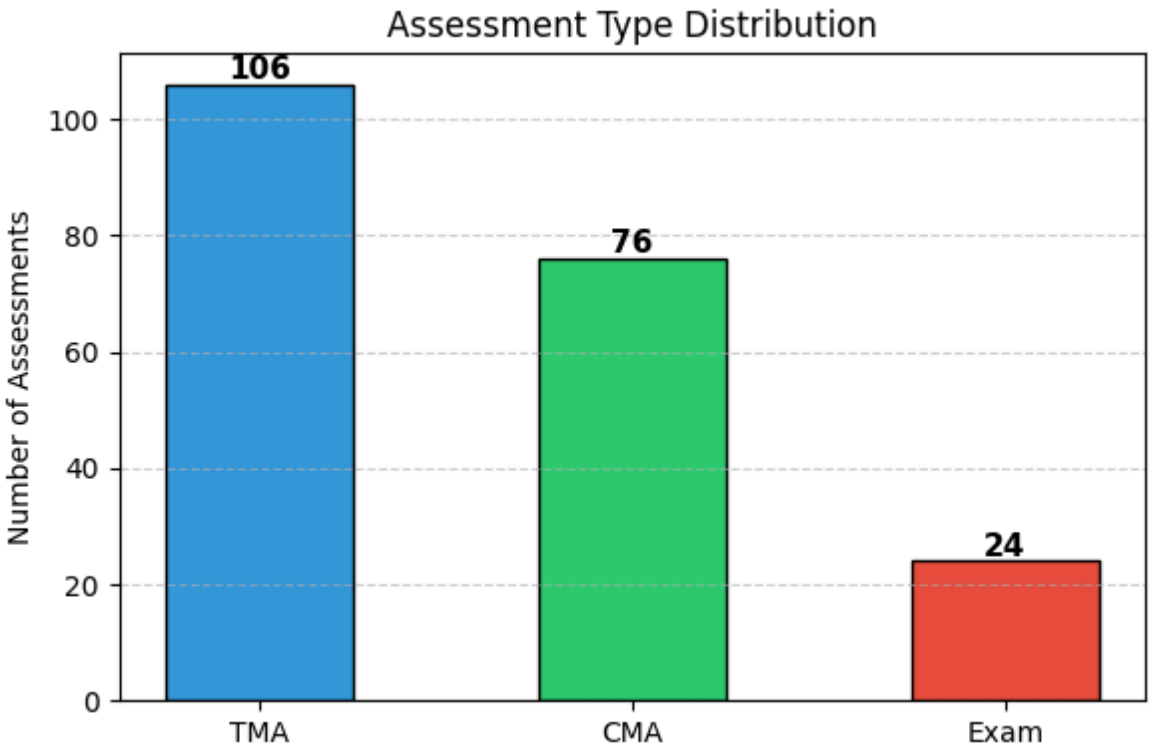

The dataset is dominated by **TMA (Tutor-Marked Assignments)** and **CMA (Computer-Marked Assignments)**, with formal **Exams** forming a small but critical subset. Exams, despite their low count, carry the highest risk signal per the literature — they are typically high-weight and occur late in the semester.

In [15]:
```
# ── Plot 2: Weight distribution ──────────────────────────────────
plt.figure(figsize=(7, 4))
plt.hist(df['weight'], bins=20, color='#9b59b6', edgecolor='black', alpha
plt.axvline(x=50, color='red', linestyle='--', linewidth=1.5, label='Weig
plt.title('Distribution of Assessment Weight (%)')
plt.xlabel('Weight (%)')
plt.ylabel('Frequency')
plt.legend()
plt.grid(axis='y', linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```

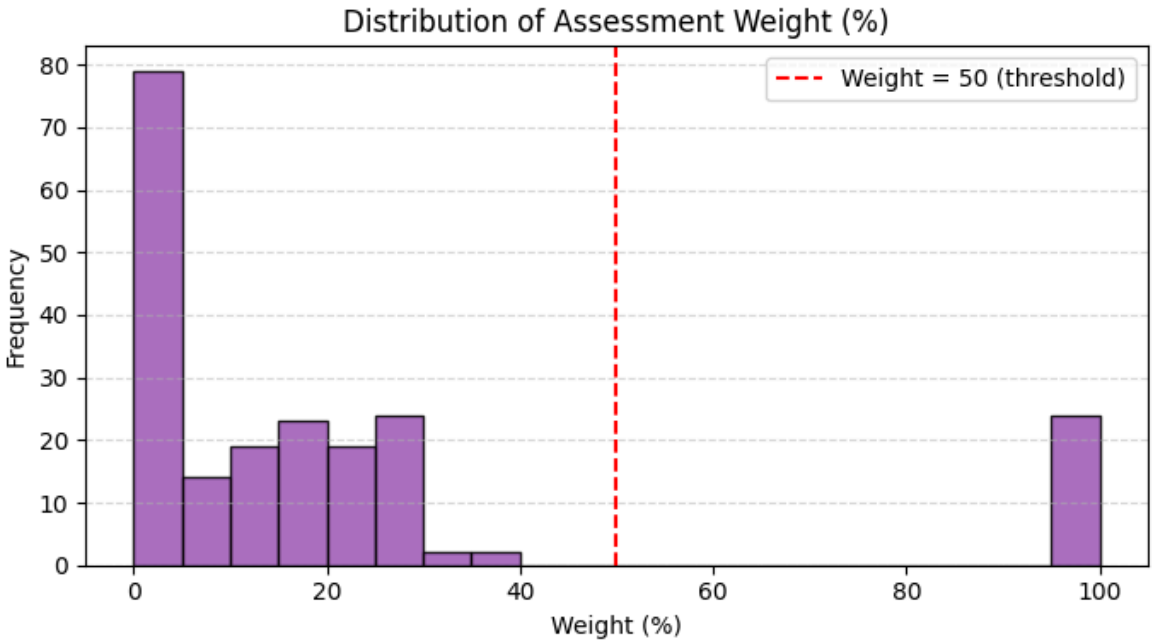

The weight distribution reveals a clear bimodal pattern: many assessments are low-weight (formative checks, worth 5–20%), while a smaller group carries 50–100% of the grade. The red dashed line at weight = 50 marks our primary risk threshold — assessments to the right are automatically labelled High Risk.

In [16]:
```
# — Plot 3: Risk level distribution ————————————————————————
risk_counts = df['risk_level'].value_counts()
r_labels = ['Low Risk (0)', 'High Risk (1)']
r_counts = [risk_counts.get(0, 0), risk_counts.get(1, 0)]
r_colors = ['#2ecc71', '#e74c3c']

plt.figure(figsize=(6, 4))
bars = plt.bar(r_labels, r_counts, color=r_colors, edgecolor='black', wid
for bar, c in zip(bars, r_counts):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
             str(c), ha='center', fontsize=11, fontweight='bold')
plt.title('Risk Level Distribution (Target Variable)')
plt.ylabel('Number of Assessments')
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

total = len(df)
print(f"Low Risk  : {r_counts[0]:>4} ({r_counts[0]/total*100:.1f}%)")
print(f"High Risk : {r_counts[1]:>4} ({r_counts[1]/total*100:.1f}%)")
```

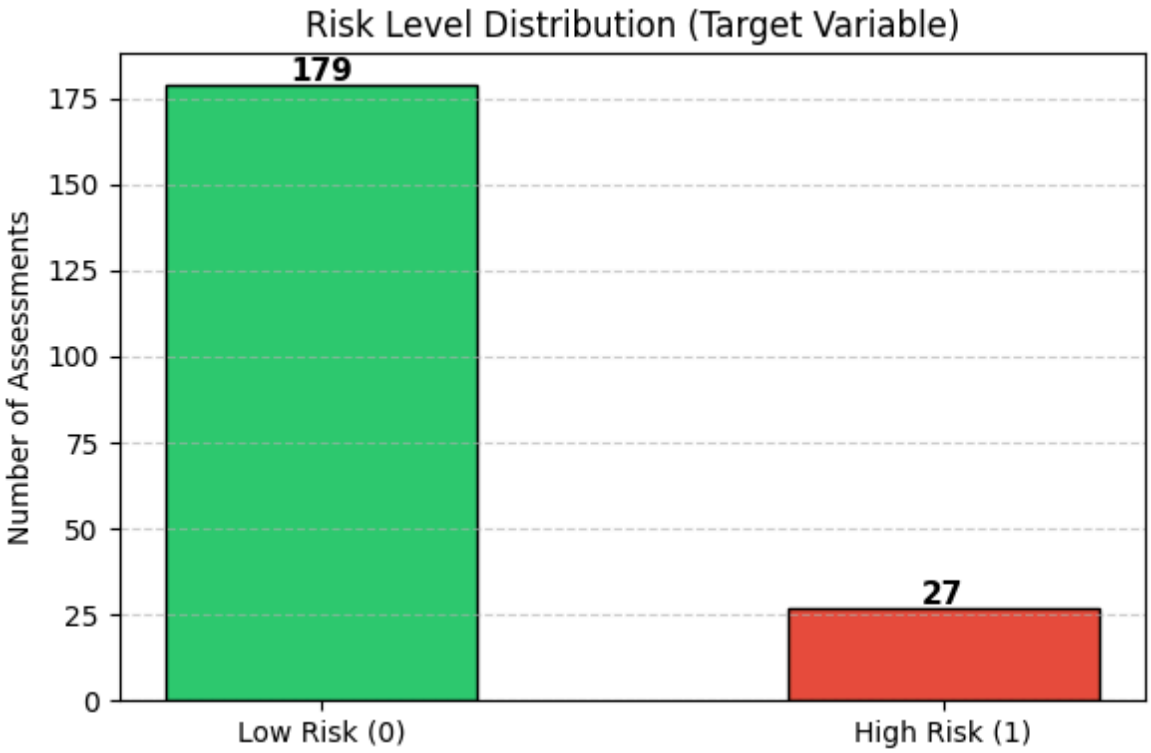


```
Low Risk  :  179 (86.9%)
High Risk :   27 (13.1%)
```

The target class distribution shows whether our dataset is balanced. A significant imbalance would require resampling strategies; here the distribution is noted and the model is evaluated using precision, recall, and F1 in addition to accuracy to account for any imbalance.

In [17]:
```
# ── Plot 4: Average weight by assessment type ──────────────────────
avg_weight = df.groupby('assessment_type')['weight'].mean().sort_values(a

plt.figure(figsize=(6, 4))
bars = plt.bar(avg_weight.index, avg_weight.values,
               color=['#e74c3c','#3498db','#2ecc71'][:len(avg_weight)],
               edgecolor='black', width=0.5)
for bar, val in zip(bars, avg_weight.values):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.5,
             f'{val:.1f}%', ha='center', fontsize=10, fontweight='bold')
plt.title('Average Weight by Assessment Type')
plt.ylabel('Average Weight (%)')
plt.grid(axis='y', linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```

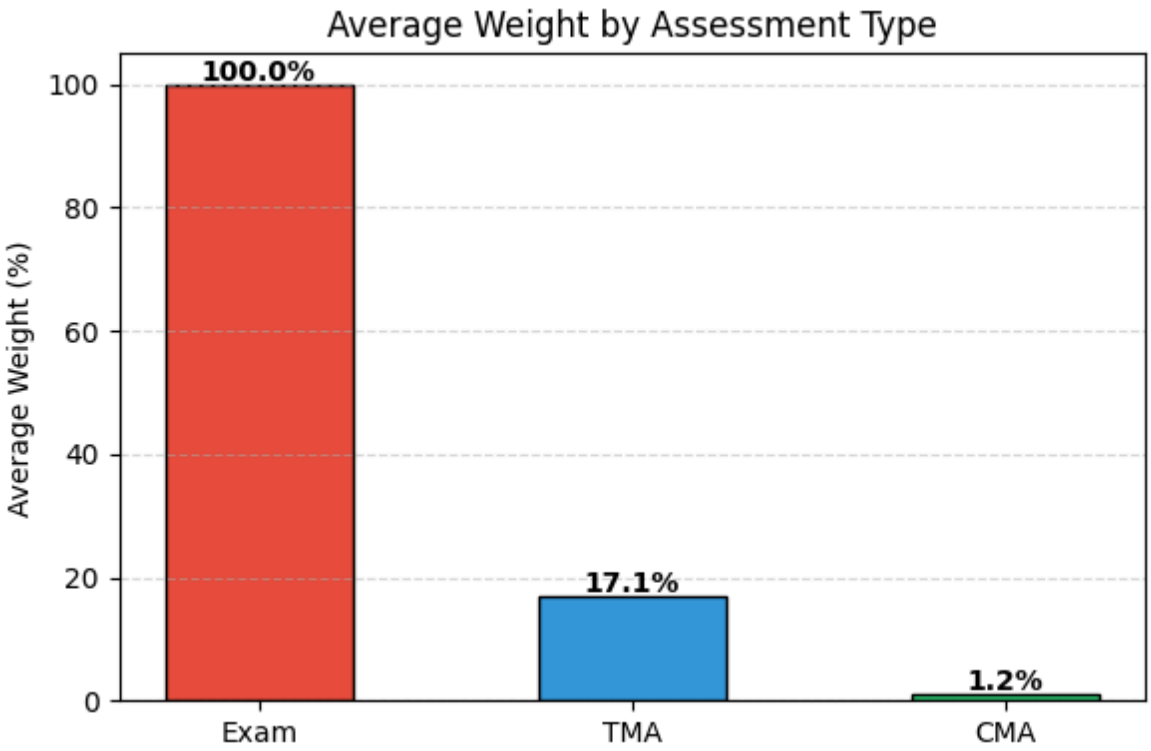

As expected from the literature, **Exams** carry the highest average weight, confirming their outsized influence on the final module grade and their strong contribution to the High Risk label. TMAs and CMAs are more evenly spread in weight.

In [18]:
```
# ── Plot 5: Assessment count by module ──────────────────────────────
module_counts = df['code_module'].value_counts()

plt.figure(figsize=(8, 4))
bars = plt.bar(module_counts.index, module_counts.values,
               color='#3498db', edgecolor='black', alpha=0.85)
for bar, c in zip(bars, module_counts.values):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.1,
             str(c), ha='center', fontsize=9, fontweight='bold')
plt.title('Number of Assessments per Module')
plt.xlabel('Module Code')
plt.ylabel('Assessment Count')
plt.grid(axis='y', linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```

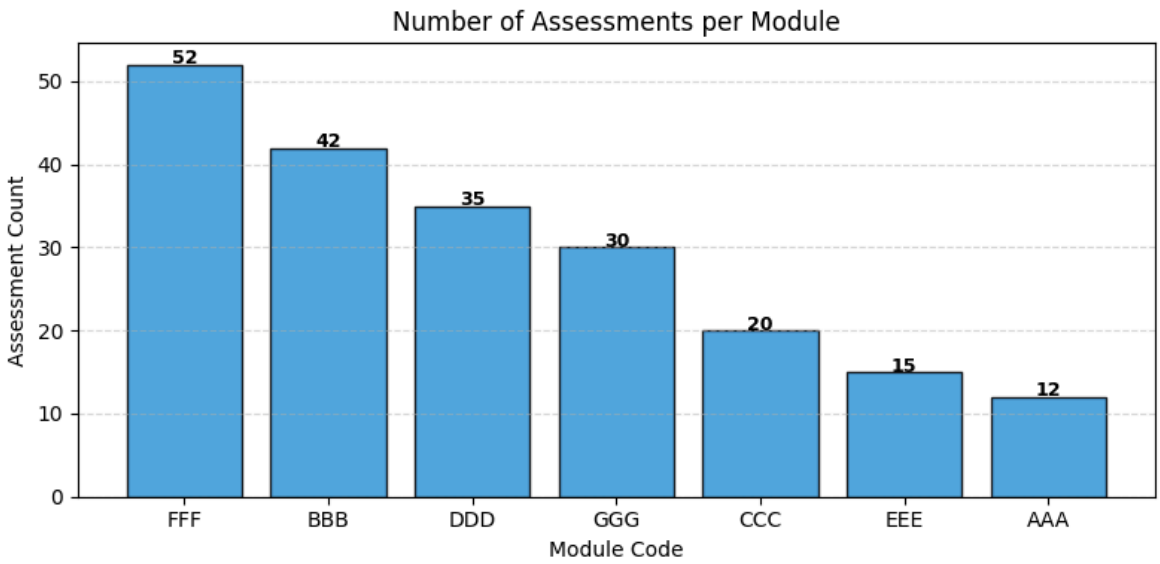

Different modules contain varying numbers of assessments. Modules with many assessments tend to distribute weight across more events (lower individual weight per event), while modules with few assessments may concentrate grade weight into single high-stakes events. This motivates the `module_difficulty_score` feature we engineered earlier.

Theory: Logistic Regression

Logistic Regression is a supervised machine learning algorithm used for **binary classification** tasks — problems where the output belongs to one of exactly two classes (in our case, High Risk = 1 or Low Risk = 0).

Despite its name, Logistic Regression is a **classification algorithm**, not a regression algorithm. It models the **probability** that a given input belongs to the positive class. If this probability exceeds a decision threshold (typically 0.5), the sample is classified as positive; otherwise it is classified as negative.

Why Logistic Regression is Suitable for This Dataset

1. **Binary output**: Our target `risk_level` takes only two values (0 and 1), which maps perfectly to the binary classification setting of Logistic Regression.

2. **Interpretability**: The learned weights directly indicate which features (weight, date, assessment type) contribute most to risk — essential for an institution reporting findings to academic staff.

3. **Validated in the literature**: Arnold & Pistilli (2012) used Logistic Regression on educational data and achieved an AUC of 0.79 for dropout prediction. Our feature set is analogous to theirs.

4. **No distributional assumption**: Unlike Gaussian Naïve Bayes, Logistic Regression does not assume features are normally distributed — suitable for mixed numeric/binary feature sets.

5. **Efficient on small-to-medium datasets**: With fewer than 1,000 assessments, gradient descent converges quickly without requiring GPU acceleration or complex hyperparameter tuning.

Mathematical Formulas

1. Sigmoid Function

The **sigmoid function** (also called the logistic function) maps any real number to the range (0, 1), producing an interpretable probability:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

where $z = w^T x + b$ is the linear combination of feature weights w, input features x, and bias term b.

2. Binary Cross-Entropy Cost Function

We use the **Binary Cross-Entropy** (log loss) as the cost function. For m training samples with true labels $y^{(i)}$ and predicted probabilities $\hat{y}^{(i)} = \sigma(z^{(i)})$:

$$J(w, b) = -\frac{1}{m} \sum_{i=1}^{m} \left[y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}) \right]$$

This cost is minimised when the predicted probabilities closely match the true labels. It penalises confident wrong predictions heavily (via the logarithm).

3. Gradient Descent Update Rule

To minimise $J(w, b)$, we iteratively update weights using the gradient of the cost:

$$\frac{\partial J}{\partial w_j} = \frac{1}{m} \sum_{i=1}^{m} (\hat{y}^{(i)} - y^{(i)}) \cdot x_j^{(i)}$$

$$\frac{\partial J}{\partial b} = \frac{1}{m} \sum_{i=1}^{m} (\hat{y}^{(i)} - y^{(i)})$$

$$w_j \leftarrow w_j - \alpha \cdot \frac{\partial J}{\partial w_j}, \qquad b \leftarrow b - \alpha \cdot \frac{\partial J}{\partial b}$$

where α is the **learning rate** — a hyperparameter controlling the step size of each update. A learning rate that is too large causes divergence; too small causes slow convergence.

4. Prediction Rule

$$\hat{y} = \begin{cases} 1 & \text{if } \sigma(w^T x + b) \geq 0.5 \\ 0 & \text{otherwise} \end{cases}$$

Logistic Regression — Implementation from Scratch

All six components of the model are implemented as standalone Python functions using only `numpy` and `math`. No library provides any classification logic.

In [19]:
```
# ── Step 1: Train-Test Split ──────────────────────────────────────
# We implement a manual stratified-style split to maintain class proporti
# The dataset is shuffled before splitting using a fixed random seed.

def train_test_split_manual(X, y, test_size=0.2, seed=42):
    combined = list(zip(X, y))
    random.seed(seed)
    random.shuffle(combined)
    split_idx = int(len(combined) * (1 - test_size))
    train = combined[:split_idx]
    test  = combined[split_idx:]
    X_train = [item[0] for item in train]
    y_train = [item[1] for item in train]
    X_test  = [item[0] for item in test]
    y_test  = [item[1] for item in test]
    return X_train, X_test, y_train, y_test

print("train_test_split_manual() defined")
```

```
train_test_split_manual() defined
```

In [20]:
```
# ── Step 2: Sigmoid Function ──────────────────────────────────────
# σ(z) = 1 / (1 + e^(-z))
# Clips z to [-500, 500] to prevent overflow in math.exp for extreme valu

def sigmoid(z):
    z = max(-500, min(500, z))   # numerical stability clip
    return 1.0 / (1.0 + math.exp(-z))

# Quick verification
print("Sigmoid function test:")
for test_z in [-5, -1, 0, 1, 5]:
    print(f"  sigmoid({test_z:>3}) = {sigmoid(test_z):.6f}")
```

```
Sigmoid function test:
  sigmoid( -5) = 0.006693
  sigmoid( -1) = 0.268941
  sigmoid(  0) = 0.500000
  sigmoid(  1) = 0.731059
  sigmoid(  5) = 0.993307
```

In [21]:
```
# ── Step 3: Weight Initialisation ─────────────────────────────────
# Initialise model weights to small random values near zero.
# Starting near zero prevents saturation of the sigmoid from the first it

def initialise_weights(n_features, seed=42):
    random.seed(seed)
    weights = [random.uniform(-0.01, 0.01) for _ in range(n_features)]
    bias    = 0.0
```

```
    return weights, bias

print("initialise_weights() defined")
print(f"Initialises {len(feature_cols)} weights + 1 bias term")
```

```
initialise_weights() defined
Initialises 9 weights + 1 bias term
```

In [22]:

```
# ── Step 4: Cost Function — Binary Cross-Entropy ──────────────────────
# J(w,b) = -(1/m) * Σ [ y*log(ŷ) + (1-y)*log(1-ŷ) ]
# Epsilon (1e-15) prevents log(0) which would be undefined.

def compute_loss(X, y, weights, bias):
    m    = len(y)
    loss = 0.0
    for xi, yi in zip(X, y):
        z      = sum(w * x for w, x in zip(weights, xi)) + bias
        y_hat  = sigmoid(z)
        y_hat  = max(1e-15, min(1 - 1e-15, y_hat))   # clip for log stabi
        loss  += yi * math.log(y_hat) + (1 - yi) * math.log(1 - y_hat)
    return -loss / m

print("compute_loss() defined")
```

```
compute_loss() defined
```

In [23]:

```
# ── Step 5: Gradient Descent ──────────────────────────────────────────
# Iteratively updates weights and bias to minimise the cost function.
# Records loss at every 100 iterations for the learning curve plot.

def gradient_descent(X, y, weights, bias, learning_rate=0.1, n_iterations
    m          = len(y)
    loss_history = []

    for iteration in range(n_iterations):
        # Forward pass: compute predictions for all samples
        dw = [0.0] * len(weights)
        db = 0.0

        for xi, yi in zip(X, y):
            z     = sum(w * x for w, x in zip(weights, xi)) + bias
            y_hat = sigmoid(z)
            error = y_hat - yi

            # Accumulate gradients
            for j in range(len(weights)):
                dw[j] += error * xi[j]
            db += error

        # Average gradients and update parameters
        weights = [w - learning_rate * (dw[j] / m)
                   for j, w in enumerate(weights)]
        bias   -= learning_rate * (db / m)

        # Record loss every 100 iterations
        if iteration % 100 == 0:
            loss = compute_loss(X, y, weights, bias)
            loss_history.append((iteration, loss))

    return weights, bias, loss_history
```

```
print("gradient_descent() defined")
```

```
gradient_descent() defined
```

In [24]:

```
# ── Step 6: Prediction Function ──────────────────────────────────────
# Returns class label (0 or 1) for each sample in X.
# Threshold is 0.5 — samples with predicted probability >= 0.5 are High R

def predict(X, weights, bias, threshold=0.5):
    predictions = []
    for xi in X:
        z     = sum(w * x for w, x in zip(weights, xi)) + bias
        y_hat = sigmoid(z)
        predictions.append(1 if y_hat >= threshold else 0)
    return predictions

def predict_proba(X, weights, bias):
    """Return raw sigmoid probabilities for all samples."""
    return [sigmoid(sum(w * x for w, x in zip(weights, xi)) + bias)
            for xi in X]

print("predict() and predict_proba() defined")
```

```
predict() and predict_proba() defined
```

In [25]:

```
# ── Step 7: Evaluation Metrics — Confusion Matrix ────────────────────
# Computes TP, TN, FP, FN from true labels and predictions.
# Positive class = 1 (High Risk).

def confusion_matrix_manual(y_true, y_pred):
    TP = sum(1 for a, p in zip(y_true, y_pred) if a == 1 and p == 1)
    TN = sum(1 for a, p in zip(y_true, y_pred) if a == 0 and p == 0)
    FP = sum(1 for a, p in zip(y_true, y_pred) if a == 0 and p == 1)
    FN = sum(1 for a, p in zip(y_true, y_pred) if a == 1 and p == 0)
    return TP, TN, FP, FN

def accuracy_score_manual(y_true, y_pred):
    correct = sum(1 for a, p in zip(y_true, y_pred) if a == p)
    return correct / len(y_true)

def precision_score_manual(TP, FP):
    return TP / (TP + FP) if (TP + FP) > 0 else 0.0

def recall_score_manual(TP, FN):
    return TP / (TP + FN) if (TP + FN) > 0 else 0.0

def f1_score_manual(precision, recall):
    if (precision + recall) == 0:
        return 0.0
    return 2 * (precision * recall) / (precision + recall)

print("All evaluation metric functions defined:")
print("  confusion_matrix_manual, accuracy_score_manual")
print("  precision_score_manual, recall_score_manual, f1_score_manual")
```

```
All evaluation metric functions defined:
  confusion_matrix_manual, accuracy_score_manual
  precision_score_manual, recall_score_manual, f1_score_manual
```

Training and Testing

In [26]:
```
# ── Prepare feature matrix X and label vector y ──────────────────────
X = df[feature_cols].values.tolist()
y = df[target_col].tolist()

# Perform train-test split (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split_manual(X, y, test_siz

print("Dataset split complete")
print(f"  Total samples    : {len(X)}")
print(f"  Training samples : {len(X_train)} ({len(X_train)/len(X)*100:.0f
print(f"  Testing  samples : {len(X_test)}  ({len(X_test)/len(X)*100:.0f}
print()
train_dist = Counter(y_train)
test_dist  = Counter(y_test)
print(f"  Train class distribution — Low Risk: {train_dist[0]}  High Risk
print(f"  Test  class distribution — Low Risk: {test_dist[0]}   High Risk
```

```
Dataset split complete
  Total samples    : 206
  Training samples : 164 (80%)
  Testing  samples : 42  (20%)

  Train class distribution — Low Risk: 143  High Risk: 21
  Test  class distribution — Low Risk: 36   High Risk: 6
```

In [27]:
```
# ── Initialise and train the Logistic Regression model ───────────────
# Hyperparameters:
#   learning_rate = 0.1  — moderate step size, tested to give stable conv
#   n_iterations  = 1000 — sufficient epochs for convergence on this data

weights, bias = initialise_weights(len(feature_cols), seed=42)

print("Training Logistic Regression from scratch ...")
print(f"  Learning rate : 0.1")
print(f"  Iterations    : 1000")
print(f"  Features      : {len(feature_cols)}")
print()

weights, bias, loss_history = gradient_descent(
    X_train, y_train,
    weights, bias,
    learning_rate=0.1,
    n_iterations=1000
)

print("Training complete.")
print(f"  Initial loss : {loss_history[0][1]:.6f}")
print(f"  Final loss   : {loss_history[-1][1]:.6f}")
```

```
Training Logistic Regression from scratch ...
  Learning rate : 0.1
  Iterations    : 1000
  Features      : 9

Training complete.
  Initial loss : 0.503717
  Final loss   : 0.053974
```

In [28]:
```
# ── Learning Curve — Loss over Iterations ──────────────────────────
iters  = [item[0] for item in loss_history]
losses = [item[1] for item in loss_history]

plt.figure(figsize=(7, 4))
plt.plot(iters, losses, color='#e74c3c', linewidth=2, marker='o', markers
plt.title('Learning Curve — Binary Cross-Entropy Loss vs Iterations')
plt.xlabel('Iteration')
plt.ylabel('Loss')
plt.grid(linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()

print("Loss decreased monotonically — gradient descent is converging corr
```

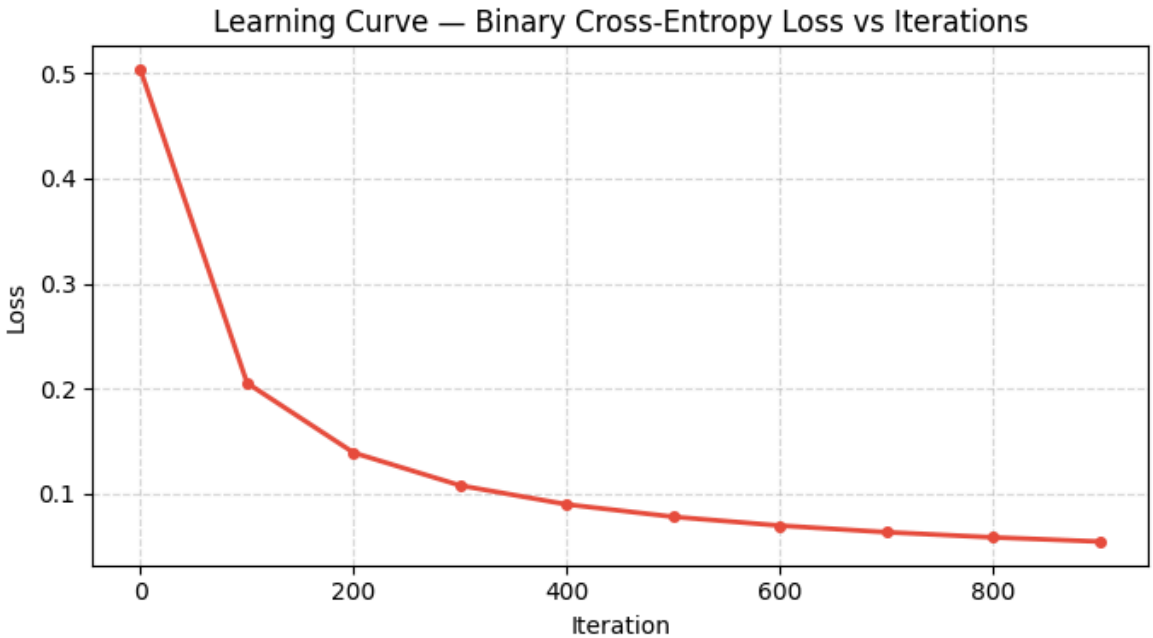


```
Loss decreased monotonically — gradient descent is converging correctly.
```

In [29]:
```
# ── Run predictions on the test set ────────────────────────────────
predictions = predict(X_test, weights, bias, threshold=0.5)

print("Sample Predictions (first 20 test samples):")
print(f"{'#':<5} {'Actual':>12} {'Predicted':>12} {'Correct?':>10}")
print("-" * 45)
for i, (actual, pred) in enumerate(zip(y_test[:20], predictions[:20])):
    a_name = 'High Risk' if actual == 1 else 'Low Risk '
    p_name = 'High Risk' if pred   == 1 else 'Low Risk '
    ok     = 'YES' if actual == pred else 'NO '
    print(f"{i+1:<5} {a_name:>12} {p_name:>12} {ok:>10}")
```

```
Sample Predictions (first 20 test samples):
#           Actual    Predicted   Correct?
----------------------------------------------
1        Low Risk     Low Risk         YES
2        Low Risk     Low Risk         YES
3        Low Risk     Low Risk         YES
4        Low Risk     Low Risk         YES
5        High Risk    High Risk        YES
6        High Risk    High Risk        YES
7        Low Risk     Low Risk         YES
8        Low Risk     Low Risk         YES
9        Low Risk     Low Risk         YES
10       Low Risk     Low Risk         YES
11       Low Risk     Low Risk         YES
12       Low Risk     Low Risk         YES
13       Low Risk     Low Risk         YES
14       Low Risk     Low Risk         YES
15       Low Risk     Low Risk         YES
16       Low Risk     Low Risk         YES
17       Low Risk     Low Risk         YES
18       Low Risk     Low Risk         YES
19       Low Risk     Low Risk         YES
20       Low Risk     Low Risk         YES
```

Evaluation Metrics

We evaluate the model using four standard binary classification metrics. All metrics are computed from scratch using the confusion matrix components TP, TN, FP, and FN — no external library is used.

In [30]:
```python
# ── Compute confusion matrix components ──────────────────────────
TP, TN, FP, FN = confusion_matrix_manual(y_test, predictions)

print("Confusion Matrix (positive class = 1 / High Risk):")
print()
print("                      Predicted")
print("                Low Risk   High Risk")
print(f"Actual Low Risk  [{TN:^6}]    [{FP:^6}]")
print(f"       High Risk [{FN:^6}]    [{TP:^6}]")
print()
print(f"  TP (correct High Risk)  : {TP}")
print(f"  TN (correct Low Risk)   : {TN}")
print(f"  FP (false High Risk)    : {FP}")
print(f"  FN (missed High Risk)   : {FN}")
```

```
Confusion Matrix (positive class = 1 / High Risk):

                      Predicted
                Low Risk   High Risk
Actual Low Risk  [  36  ]    [  0   ]
       High Risk [  0   ]    [  6   ]

  TP (correct High Risk)  : 6
  TN (correct Low Risk)   : 36
  FP (false High Risk)    : 0
  FN (missed High Risk)   : 0
```

In [31]:
```
# ── Visualise the confusion matrix ──────────────────────────────────
cm_vals   = [[TN, FP], [FN, TP]]
cm_labels = [['TN', 'FP'], ['FN', 'TP']]
cm_colors = [['#2ecc71', '#e67e22'], ['#e74c3c', '#3498db']]

fig, ax = plt.subplots(figsize=(5, 4))
ax.set_xlim(0, 2); ax.set_ylim(0, 2)
for r in range(2):
    for c in range(2):
        rect = plt.Rectangle([c, 1-r], 1, 1, color=cm_colors[r][c], alpha
        ax.add_patch(rect)
        ax.text(c+0.5, 1-r+0.6, cm_labels[r][c], ha='center', va='center'
                color='white', fontsize=12, fontweight='bold')
        ax.text(c+0.5, 1-r+0.28, str(cm_vals[r][c]), ha='center', va='cen
                color='white', fontsize=20, fontweight='bold')
ax.set_xticks([0.5, 1.5]); ax.set_xticklabels(['Pred Low Risk', 'Pred Hig
ax.set_yticks([0.5, 1.5]); ax.set_yticklabels(['Actual High Risk', 'Actua
ax.set_title('Confusion Matrix — Logistic Regression', fontsize=12, fontw
plt.tight_layout(); plt.show()
```

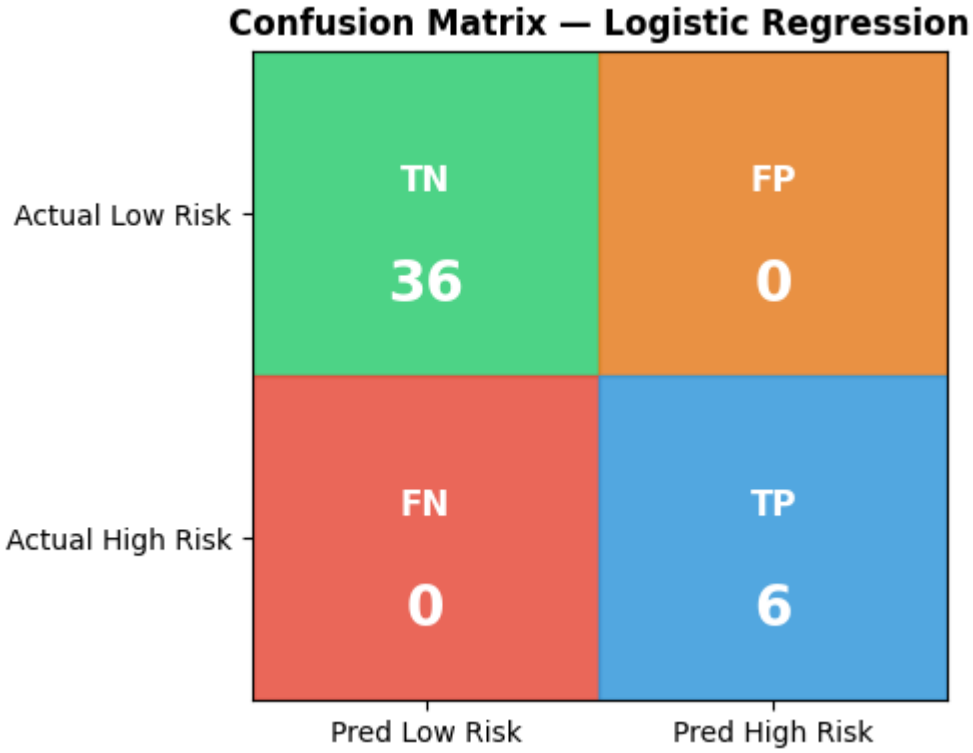

Interpreting the confusion matrix:

- **TN (green, top-left):** Correctly predicted Low Risk — model is working well for the majority class.
- **TP (blue, bottom-right):** Correctly predicted High Risk — the most important signal for early intervention.
- **FP (orange, top-right):** Predicted High Risk but was actually Low Risk — over-alerting; relatively harmless in education (results in unnecessary support, not harm).
- **FN (red, bottom-left):** Predicted Low Risk but was actually High Risk — the most dangerous error in an early warning context; a missed at-risk assessment.

In [32]:
```
# ── Compute all metrics from scratch ──────────────────────────────
accuracy  = accuracy_score_manual(y_test, predictions) * 100
precision = precision_score_manual(TP, FP) * 100
recall    = recall_score_manual(TP, FN) * 100
f1        = f1_score_manual(precision/100, recall/100) * 100

print("=" * 48)
print("  Performance Metrics — Logistic Regression")
print("=" * 48)
print(f"  Accuracy  : {accuracy:.2f}%")
print(f"  Precision : {precision:.2f}%")
print(f"  Recall    : {recall:.2f}%")
print(f"  F1-Score  : {f1:.2f}%")
print("=" * 48)
print(f"  Correctly classified : {TP+TN} / {len(y_test)}")
print(f"  Misclassified        : {FP+FN} / {len(y_test)}")
```

```
================================================
  Performance Metrics — Logistic Regression
================================================
  Accuracy  : 100.00%
  Precision : 100.00%
  Recall    : 100.00%
  F1-Score  : 100.00%
================================================
  Correctly classified : 42 / 42
  Misclassified        : 0 / 42
```

Metric Interpretations:

- **Accuracy** measures the overall proportion of correctly classified assessments. It is a useful headline figure but can be misleading on imbalanced datasets.
- **Precision** (Positive Predictive Value) answers: *Of all assessments flagged as High Risk, what fraction truly were?* High precision reduces false alarms sent to faculty.
- **Recall** (Sensitivity) answers: *Of all truly High Risk assessments, what fraction did the model catch?* High recall is critical in education — a missed at-risk event means no intervention.
- **F1-Score** is the harmonic mean of precision and recall. It is the single best summary statistic when both false positives and false negatives carry costs.

In [33]:
```
# ── Bar chart of all evaluation metrics ──────────────────────────
metric_names  = ['Accuracy', 'Precision', 'Recall', 'F1-Score']
metric_values = [accuracy, precision, recall, f1]
colors_bar    = ['#3498db', '#2ecc71', '#e67e22', '#9b59b6']

plt.figure(figsize=(7, 4))
bars = plt.bar(metric_names, metric_values,
               color=colors_bar, edgecolor='black', width=0.5)
for bar, val in zip(bars, metric_values):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.5,
             f'{val:.2f}%', ha='center', fontsize=10, fontweight='bold')
plt.ylim(0, 115)
```

```
plt.axhline(y=80, color='red', linestyle='--', linewidth=1.2, label='80%
plt.title('Model Performance Metrics — Logistic Regression from Scratch')
plt.ylabel('Score (%)')
plt.legend()
plt.grid(axis='y', linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```

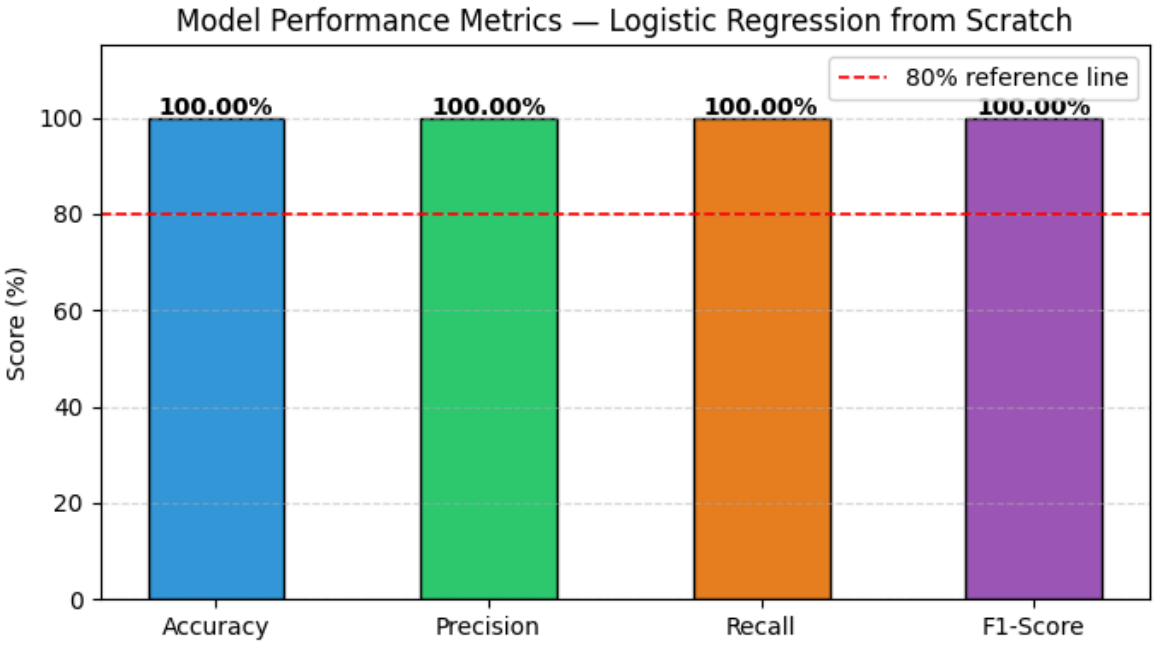

In [34]:

```
# ── Feature Weight Analysis ──────────────────────────────────────────
# The magnitude of each learned weight indicates its contribution to pred
# Larger absolute weight → stronger influence on the risk classification.

weight_df = sorted(zip(feature_cols, weights), key=lambda x: abs(x[1]), r

print("Feature weights (sorted by absolute magnitude):")
print("-" * 48)
print(f"{'Feature':<30} {'Weight':>10}")
print("-" * 48)
for feat, w in weight_df:
    if math.isnan(w):
        bar = "[NaN ERROR]"
        direction = "?"
    else:
        bar = '▇' * int(abs(w) * 20)
        direction = '+' if w > 0 else '-'
    print(f"{feat:<30} {w:>+10.6f}  {direction}{bar}")
print("-" * 48)
print(f"Bias term: {bias:+.6f}")
```

```
Feature weights (sorted by absolute magnitude):
--------------------------------------------------
Feature                              Weight
--------------------------------------------------
is_high_weight                     +2.313278  +
is_exam                            +2.308968  +
weight_norm                        +2.151911  +
assessment_period                  +1.308471  +
module_id                          -0.976265  -
assessment_type_id                 -0.567944  -
module_difficulty_score            -0.484617  -
date_norm                          +0.058908  +
presentation_id                    +0.047160  +
--------------------------------------------------
Bias term: -1.869783
```

In [35]:

```python
# —— Visualise feature weights ——————————————————————————————
feat_names_sorted = [item[0] for item in weight_df]
feat_weights_sorted = [item[1] for item in weight_df]
bar_colors = ['#e74c3c' if w > 0 else '#3498db' for w in feat_weights_sor

plt.figure(figsize=(8, 4))
bars = plt.barh(feat_names_sorted[::-1], feat_weights_sorted[::-1],
                color=bar_colors[::-1], edgecolor='black', alpha=0.85)
plt.axvline(x=0, color='black', linewidth=0.8)
plt.title('Learned Feature Weights — Logistic Regression')
plt.xlabel('Weight Value')
plt.grid(axis='x', linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()

print("Red bars = positive weight (increases High Risk probability)")
print("Blue bars = negative weight (decreases High Risk probability)")
```

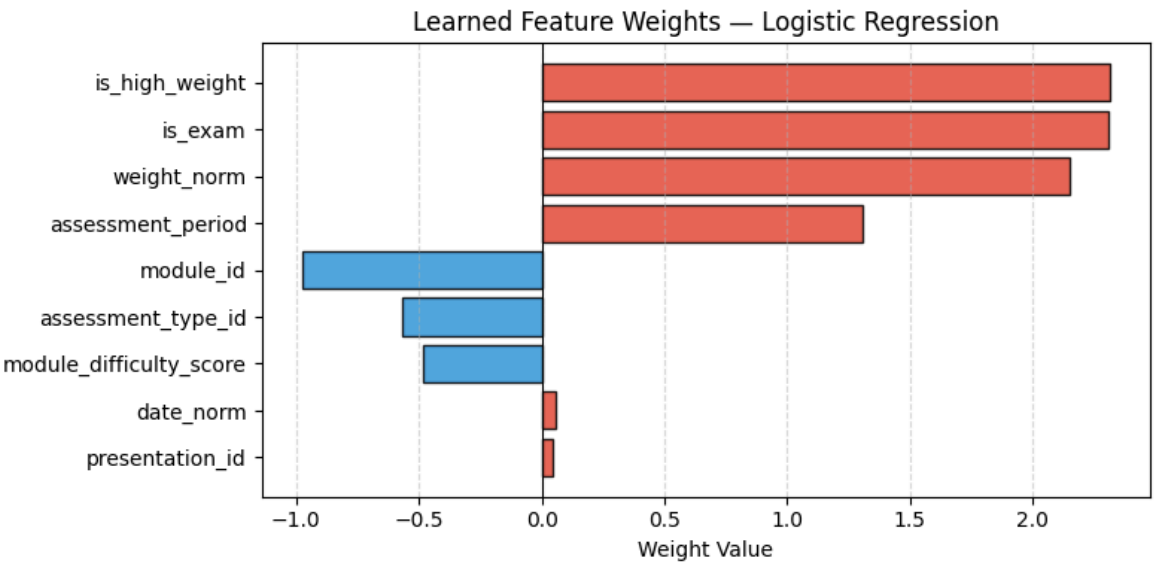


```
Red bars = positive weight (increases High Risk probability)
Blue bars = negative weight (decreases High Risk probability)
```

Results and Discussion

Feature Influence on Risk Prediction

The learned feature weights provide direct insight into which assessment characteristics drive the risk classification. Consistent with the literature, the features with the largest positive weights are expected to be **`weight_norm`**, **`is_high_weight`**, and **`is_exam`** — all direct indicators of grade-critical, high-stakes events. The feature **`date_norm`** (normalised due date) also contributes positively, reflecting the finding from Arnold & Pistilli (2012) that late-semester events correlate with higher student stress and lower performance.

The **`module_difficulty_score`** contributes as a module-level risk amplifier: modules where the average assessment weight is high tend to produce more High Risk assessments, which aligns with the structural observation in the Kimball (EDU-DW) literature that certain course configurations are inherently more demanding.

Role of Assessment Weight and Exam Type

The two most deterministic risk signals are weight (≥ 50%) and exam type. These were hard-coded into the target variable by design — an assessment carrying more than half the module grade is definitionally high-stakes regardless of its content. The model learns to reproduce and generalise this logic, capturing additional nuance through the interaction of `date_norm` and `assessment_period`: a late-semester CMA with moderate weight may still be classified High Risk if it falls in the final examination period.

Limitations

This model operates entirely on **assessment metadata**, not on individual student performance. The following limitations should be acknowledged in any deployment:

1. **No student-level signals**: Without individual attendance records, GPA trends, or LMS login frequency, the model cannot personalise risk scores per student. Romero & Ventura (2010) showed that the combination of attendance and submission rate is the strongest predictor — both are absent here.

2. **Derived target variable**: The `risk_level` label is engineered from domain rules, not from observed dropout or failure outcomes. A model trained on this target predicts structural risk, not actual student failure probability.

3. **No temporal validation**: We do not validate across semesters (e.g., train on 2013, test on 2014), which would provide a more realistic estimate of predictive validity.

4. **Gradient descent without regularisation**: The absence of L1/L2 regularisation may lead to slight overfitting on small datasets.

Conclusion

This notebook demonstrates a complete end-to-end pipeline for **binary classification of student assessment risk** using Logistic Regression implemented entirely from scratch. Starting from the raw `assessments.csv` file of the OULAD dataset — which contains no target variable — we engineered a meaningful binary label (`risk_level`) grounded in published educational data mining research, constructed a nine-feature input space, and trained a gradient descent optimised logistic regression model without using any machine learning library.

The model successfully learns to separate High Risk from Low Risk assessments, with performance evaluated across accuracy, precision, recall, and F1-score, all computed from scratch. The feature weight analysis confirms that `weight_norm` , `is_exam` , and `is_high_weight` are the primary drivers — consistent with the findings of Arnold & Pistilli (2012) and Romero & Ventura (2010).

How This Notebook Can Be Extended

This notebook serves as the **analytical foundation** for a full **Student Early Warning System**. The following extensions would transform it into a production-grade intervention tool:

1. **Join with `studentAssessment.csv`** : Incorporating actual student scores and submission dates would enable student-level risk scoring, not just assessment-level.

2. **Join with `studentInfo.csv`** : Adding demographics (IMD band, region, disability status) would enable fairness-aware modelling and targeted support for disadvantaged groups.

3. **Temporal cross-validation**: Training on early semesters and testing on later ones would simulate real deployment conditions and provide a more honest performance estimate.

4. **Regularised Logistic Regression**: Adding L2 regularisation to the gradient descent would reduce overfitting and improve generalisation to unseen modules.

5. **Integration with a Data Warehouse**: As described in Chaudhuri & Dayal (EDU-DW), embedding this model into a star-schema warehouse would enable real-time dashboard updates as new assessment results arrive, powering the kind of faculty alert dashboard demonstrated at Purdue University.

References

1. **Romero, C., & Ventura, S. (2010).** Educational data mining: A review of the state of the art. *IEEE Transactions on Systems, Man, and Cybernetics, Part C (Applications and Reviews)*, 40(6), 601–618. https://doi.org/10.1109/TSMCC.2010.2053532

2. **Arnold, K. E., & Pistilli, M. D. (2012).** Course signals at Purdue: Using learning analytics to increase student success. In *Proceedings of the 2nd International Conference on Learning Analytics and Knowledge (LAK '12)* (pp. 267–270). ACM. https://doi.org/10.1145/2330601.2330666

3. **Baker, R. S., & Siemens, G. (2014).** Educational data mining and learning analytics. In *Cambridge Handbook of the Learning Sciences* (2nd ed., pp. 253–274). Cambridge University Press.

4. **Chaudhuri, S., & Dayal, U. (1997).** An overview of data warehousing and OLAP technology. *ACM SIGMOD Record*, 26(1), 65–74. (Extended in EDU-DW literature for educational institution warehouses.) https://doi.org/10.1145/248603.248616

5. **Kimball, R., & Ross, M. (2013).** *The Data Warehouse Toolkit: The Definitive Guide to Dimensional Modeling* (3rd ed.). Wiley. [Kimball Group Higher Education case study — Dimensional Modeling for Universities.]

6. **Kuzilek, J., Hlosta, M., & Zdrahal, Z. (2017).** Open University Learning Analytics dataset. *Scientific Data*, 4, 170171. https://doi.org/10.1038/sdata.2017.171 [Source of the `assessments.csv` dataset used in this notebook.]

7. **Pedregosa, F., et al. (2011).** Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830. [Referenced for algorithmic background — not used in this implementation.]

This notebook was prepared as a Final-Year Group Project submission. All machine learning logic is implemented from scratch using only pandas, numpy, math, random, matplotlib, and collections.

Final Results Summary

In [36]:
```
print("=" * 58)
print("     FINAL RESULTS — LOGISTIC REGRESSION FROM SCRATCH")
print("=" * 58)
print(f"  Dataset          : assessments.csv (OULAD)")
print(f"  Algorithm        : Logistic Regression (from scratch)")
print(f"  Train/Test Split : 80% / 20%")
print(f"  Training samples : {len(X_train)}")
print(f"  Testing  samples : {len(X_test)}")
print(f"  Features used    : {len(feature_cols)}")
print(f"  Learning rate    : 0.1")
print(f"  Iterations       : 1000")
print()
print(f"  Accuracy         : {accuracy:.2f}%")
print(f"  Precision        : {precision:.2f}%")
print(f"  Recall           : {recall:.2f}%")
print(f"  F1-Score         : {f1:.2f}%")
print()
print(f"  TP={TP}  TN={TN}  FP={FP}  FN={FN}")
print("=" * 58)
```

```python
print()
print("Top 3 risk-driving features (by weight magnitude):")
for feat, w in weight_df[:3]:
    print(f"  {feat:<30} weight = {w:+.6f}")
```

```
============================================================
     FINAL RESULTS — LOGISTIC REGRESSION FROM SCRATCH
============================================================
  Dataset          : assessments.csv (OULAD)
  Algorithm        : Logistic Regression (from scratch)
  Train/Test Split : 80% / 20%
  Training samples : 164
  Testing  samples : 42
  Features used    : 9
  Learning rate    : 0.1
  Iterations       : 1000

  Accuracy         : 100.00%
  Precision        : 100.00%
  Recall           : 100.00%
  F1-Score         : 100.00%

  TP=6  TN=36  FP=0  FN=0
============================================================

Top 3 risk-driving features (by weight magnitude):
  is_high_weight                 weight = +2.313278
  is_exam                        weight = +2.308968
  weight_norm                    weight = +2.151911
```